

Comparison	Removed (M) Q / MLP	Mean Δ minADE@6 95% CI (m)	Median Δ minADE@6 95% CI (m)
Q25 – MLP25 Ada · equal unit fraction	75.5 / 456.5	+0.0316 [−0.0356, +0.0761]	+0.0205 [+0.0103, +0.0380]
Q25 – MLP4.13 Ada · matched parameters	75.5 / 75.4	+0.0317 [−0.0355, +0.0761]	+0.0217 [+0.0103, +0.0375]
Q50 – MLP50 Blackwell · equal unit fraction	151.0 / 913.0	+0.2459 [+0.1694, +0.3168]	+0.1171 [+0.0848, +0.1815]

Dense VLM; val500. Δ = Q-head pruning minus MLP-channel pruning; positive means Q pruning has higher error.
Each pair uses the same GPU architecture and identical generated CoC. Q25 removes 4/16 heads; Q50 removes 8/16.
Parameter-matched MLP removes 341/8256 channels. Mean and median are distinct estimands.