

Method	Removed (%)	Val500		Test500		OOD-val262		Closed-loop · 150 scenes × 2 rollouts				
		minADE@6 ↓ (m)	minFDE@6 ↓ (m)	minADE@6 ↓ (m)	minFDE@6 ↓ (m)	minADE@6 ↓ (m)	minFDE@6 ↓ (m)	Scene score ↑	Collision@fault ↓ (%)	Offroad ↓ (%)	Progress ↑ (ratio)	LingoQA ↑ Acc. (%)
Dense	0.0	0.8236	2.2050	0.8417	2.2673	0.9995	2.4959	0.7496	4.0	7.0	0.7467	73.2
Wanda	24.0	2.7005	7.5139	2.9754	8.2687	2.7898	6.8840	0.6212	14.7	9.3	0.6922	9.2
LLM-Pruner	24.0	1.2794	3.3627	1.4513	3.8885	1.6622	4.1534	0.8105	5.7	3.7	0.8108	44.8
Týr-the-Pruner	24.0	0.8457	2.1991	0.9503	2.5702	1.1765	2.9896	0.7864	3.0	8.7	0.7943	34.2
Traj	24.0	0.9503	2.4924	1.0076	2.7203	1.1782	2.9597	0.7830	3.3	8.7	0.8049	37.0
CoC	24.0	1.5577	4.3241	1.5837	4.4996	1.5460	4.0096	0.6604	7.3	9.0	0.6849	30.2
Dual	24.0	0.8904	2.3304	0.9498	2.5459	1.1192	2.8514	0.8283	2.3	6.7	0.8276	68.8

Clip means of minADE@6 / minFDE@6 (m); each minimum is taken independently over the first 6 samples. Horizon: 6.4 s.
150 scenes × 2 rollouts. Collision@fault / offroad: affected rollouts (%). Progress: mean progress_clipped_rel (ratio).
Lingo-Judge accuracy on 500 questions; native VQA, unprompted.
LLM-Pruner: CoC / param_first, layers 4–33. Týr: searched allocation + reconstruction.